

Optical flow

Marko Medved

I. INTRODUCTION

In this assignment, the mean shift method was first implemented and then used to develop a tracker, which was tested on the VOT14 dataset. The model's performance was evaluated using different parameters. The tracker was further improved by introducing additional weights for the histogram, with the weights selected based on the tracked object's background.

II. EXPERIMENTS

A. Mean shift mode seeking on the given example

The Mean Shift method was tested using several different parameters to assess its performance. Firstly, the kernel size, which determines the region over which the probabilities for making a step are calculated, was varied. Additionally, different starting positions on the density function were considered to evaluate how the initial location affects the results. Lastly, three convergence criteria were tested to determine when the iteration should stop: the first, based on step size, stops the iteration when the calculated movement is smaller than one pixel; the second, using Euclidean distance, halts the process when the distance between consecutive positions is less than 2 units; and the third, based on change in probability, terminates the iteration when the change in the probability distribution is minimal, indicating convergence.

In Figure 1, we observe how different combinations of these parameters affect convergence. A larger kernel size leads to faster convergence, as the calculated steps are generally larger. However, when the kernel size is too small, the algorithm fails to converge to the local maximum because the calculated steps become smaller than a pixel. On the other hand, using an excessively large kernel can also prevent convergence to the local maximum, as illustrated by the example with the largest kernel size.

The starting position is also crucial, as it directly influences the algorithm's behavior. If the starting point is located where there is little to no probability in the surrounding area, the algorithm will fail to make any steps. Additionally, for functions with multiple local maxima, the starting position determines which maximum the algorithm will converge to.

Lastly, among the different types of convergence criteria, the sub-pixel step method is generally the most consistent, although it tends to be the slowest. It's also worth noting that the "small change in probability" criterion may not always be the best choice, especially if the probabilities change insignificantly at the start, as it can lead to premature convergence.

B. Mean shift mode seeking on custom examples

The algorithm was then tested on three custom functions, shown in Figure 2. In these tests, the sub-pixel convergence criterion was used, and the kernel size was set to 25. For each function, four different starting positions were evaluated. The first function is a Gaussian distribution, where the algorithm consistently converges to the global maximum, provided the starting point is not in a region with near-zero probability. The second example is the Laplacian of Gaussian (the Mexican hat function), which has multiple local maxima. Here, the




Figure 1. Comparison of the mean shift method with different kernel sizes, convergence criteria and starting positions

convergence outcome strongly depends on the starting position. The final example is a small section of the Julian Alps (around 46° North and 14° East). Due to the presence of many local maxima, the starting position determines which peak the algorithm will ascend to.




Figure 2. Comparison of the mean shift method on different functions

C. Basic tracker implementation

The tracker was implemented and tested on the entire VOT 2014 dataset. The number of failures and the tracking speed (measured on my personal laptop) are summarized in Table I. Overall, the tracker failed a total of 33 times across the dataset. While the tracker performs reliably in most cases, it struggles more with certain sequences, such as hand2, fish1, torus, and tunnel. The reasons behind these failures will be explored in detail in the next section.

D. Failure cases discussion

The cases where the tracker experienced the most failures are shown in Figure 3. In the fish1 and hand2 sequences, the failures occurred because the color distributions of the bounding box being tracked were very similar to those of the ground truth

Table I
BASIC TRACKER RESULTS

Sequence	Failures	Speed	Sequence	Failures	Speed
ball	1	1823 FPS	david	1	1042 FPS
basketball	0	510 FPS	diving	0	1328 FPS
bicycle	1	1047 FPS	drunk	1	700 FPS
bolt	2	1425 FPS	fernando	2	358 FPS
car	0	2017 FPS	fish1	3	1571 FPS
fish2	1	1224 FPS	motocross	2	504 FPS
gymnastics	0	1667 FPS	polarbear	0	2020 FPS
hand1	2	923 FPS	skating	1	1098 FPS
hand2	5	2610 FPS	sphere	0	1285 FPS
jogging	1	1704 FPS	sunshade	0	928 FPS
surfing	0	3179 FPS	torus	3	1170 FPS
trellis	2	1449 FPS	tunnel	4	2684 FPS
woman	1	2525 FPS			

bounding box, making it difficult for the tracker to detect movement. While the failures in the other two sequences, torus and tunnel, are less obvious, they can be interpreted in a similar way, with the added complexity of scale changes in both videos. In the torus sequence, as the object rotated around its axis, the color distribution of the tracked patch changed significantly, leading to tracking failure. In the tunnel sequence, when the motorbike moved further away from the camera, its scale decreased, causing the bounding box to become too large and inaccurate.

These failure cases suggest two potential improvements for the tracker. First, the issue with scale changes could be mitigated by trying different template scales and selecting the one that provides the most similarity. Second, the tracker could be improved by addressing the influence of the background, which is a factor that will be discussed in more detail in the final section.




Figure 3. Mean shift tracker failure cases

E. Parameter tuning

Four key parameters were considered: σ (kernel size), the number of bins in the histograms, the number of iterations per frame, and α (the update rate of the template). The algorithm was tested on all combinations of 4 different bin counts, 4 different kernel sizes, 3 iteration values, and 4 update rates.

To determine the optimal number of iterations and bins, the results were aggregated based on these two parameters,

as shown in Figure 4. The data indicates that using only one iteration significantly reduces performance, whereas 10 and 20 iterations produce similar results. The average FPS for 10 and 20 iterations was 1482 and 1484, respectively, showing negligible difference in processing speed.

Regarding the number of bins, it is clear that 8 or 16 bins provide the best performance. Although 8 bins slightly underperform compared to 16, they offer a noticeable increase in average FPS. However, on my personal laptop, the average FPS for 16 bins reached 2151, making it the preferred choice unless running on a computationally limited device.

Based on these findings, the optimal settings are 20 iterations and 16 bins, as this combination minimizes failures while maintaining high processing speed. These values will be used for further evaluation.

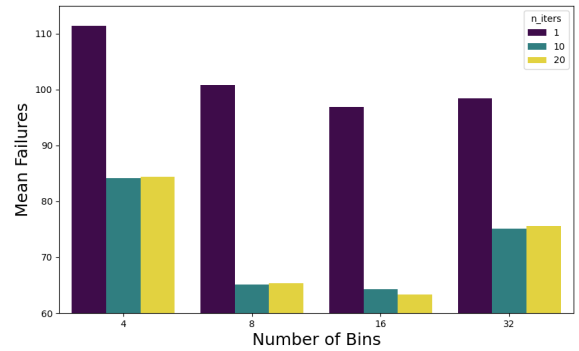


Figure 4. Performance of the means shift tracker with different numbers of bins and numbers of iterations

To determine the best values for σ and α , we refer to Figure 5. Since the FPS remains relatively consistent across different values of these parameters, the focus is on minimizing the number of failures. The results suggest that the optimal value for σ is clear. However, the choice of α may be more problem-specific, as the average performance is quite similar across the entire dataset. For scenarios where the target is expected to change frequently, $\alpha = 0.01$ is likely the better choice, whereas for more stable targets, $\alpha = 0$ would be more suitable.

F. Feature selection by accounting for the background in different color spaces

To address the background influence, the method applied smaller weights to the regions in the histogram where the background color predominated. The tracker was also evaluated using five different color spaces. Additionally, another parameter was introduced: the size of the background area. Since the optimal value for alpha was not definitive in the previous evaluation, two potential alpha values (0 and 0.01) were also considered in this experiment.

III. CONCLUSION

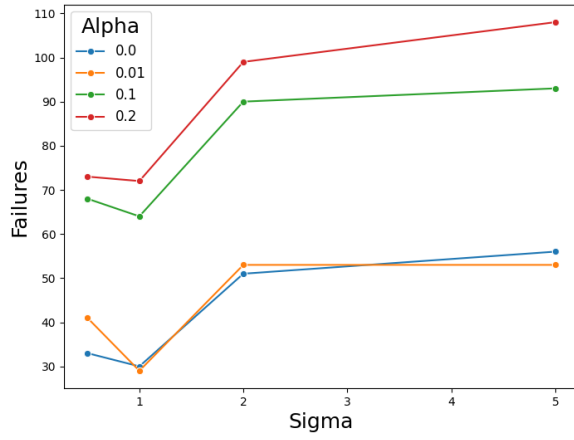


Figure 5. Performance of the means shift tracker with different kernel sizes and update rate parameters